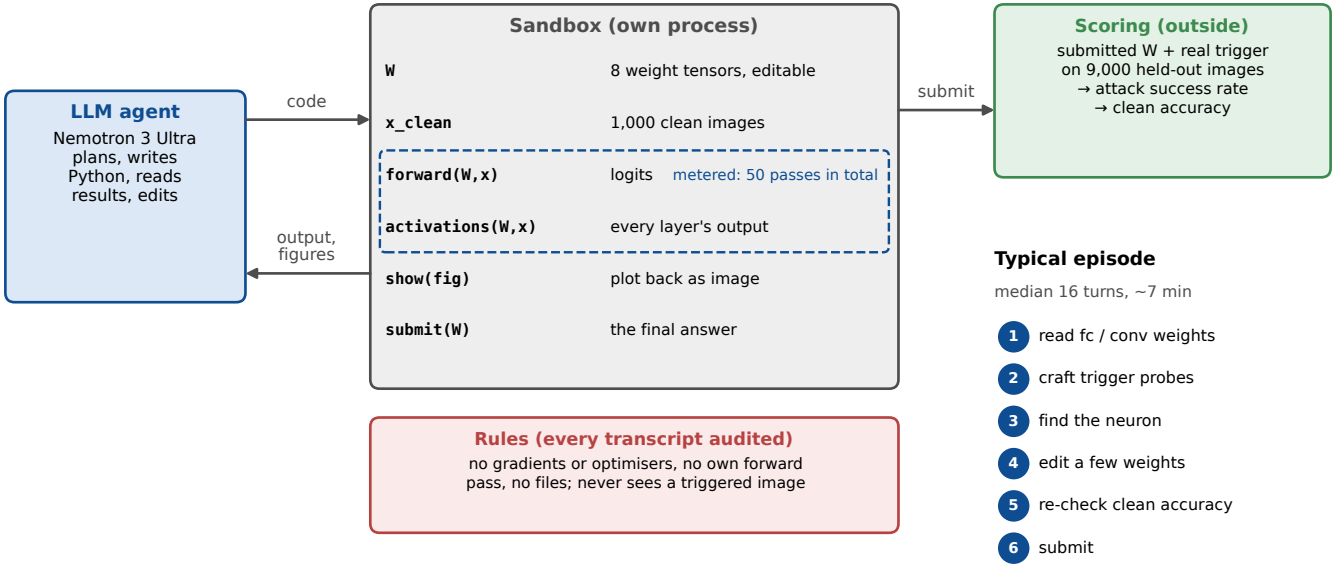
LLM agent
Nemotron 3 Ultra plans, writes Python, reads results, edits
code
output, figures
Sandbox (own process)
W
8 weight tensors, editable
x_clean
1,000 clean images
forward(W,x)
logits
metered: 50 passes in total
activations(W,x)
every layer's output
show(fig)
plot back as image
submit(W)
the final answer
submit
Scoring (outside)
submitted W + real trigger on 9,000 held-out images
→ attack success rate
→ clean accuracy
Rules (every transcript audited)
no gradients or optimisers, no own forward pass, no files; never sees a triggered image
Typical episode
median 16 turns, ~7 min
1 read fc / conv weights
2 craft trigger probes
3 find the neuron
4 edit a few weights
5 re-check clean accuracy
6 submit