§	Experiment	Question	Answer
6.1	**Zoo**	Can we plant strong backdoors that accuracy does not reveal?	Yes: mean ASR 96.7%, clean accuracy about 70% in both groups
6.2	**Detection**	Can an LLM tell from the weights alone whether a network is backdoored?	Barely: AUROC 0.61-0.65, against 0.87-0.94 for the baselines
6.3-6.4	**Removal**	Can an LLM agent remove the backdoor by editing weights, without gradients?	Partly: 6/20 removed with ~6 weights; mean ASR 51.9% vs 42.8% (NC unlearn)
6.5	**Ablations**	Does the agent succeed by reading the weights or by probing the net?	Mostly probing; reading adds 10 pt (3 runs, p = 0.009)
6.6	**Edit anatomy**	Do the edits hit the neuron that carries the backdoor?	Necessary: 0 of 69 edits that miss the backdoor channel remove it
6.7	**Repeat runs**	Does one run tell us how well the agent does on a model?	No: 11 backdoors removed in some run, only 1 in all three
6.8	**Best of 3**	Can a defender pick the good run without knowing the trigger?	Partly: picking helps; 'beats NC unlearn' only after the fact
6.9	**Pre-registered confirmation**	On 20 fresh networks, does the picked agent beat NC unlearn?	No (p = 0.69); picking itself confirmed (p = 0.0003)

supported · partly / exploratory · not supported · measured