Reading Backdoors from Weights: Can Language-Model Agents Detect and Remove Neural Trojans by Inspecting and Editing Parameters?

Yukteswar

November 14, 2026

Abstract

A backdoored neural network behaves normally on ordinary inputs but obeys a hidden trigger. We ask whether a large language model (LLM) can audit such a network the way a human expert might: by *reading* its weights to decide whether it is backdoored, and by *editing* its weights to remove the backdoor. To answer this with known ground truth, we train a zoo of 120 small CIFAR-10 convolutional networks (24,234 parameters each), half of them poisoned with BadNets-style patch triggers, and compare openly available LLMs (NVIDIA Nemotron 3 Ultra and Nemotron 3.5 Super VL) against seven conventional defences under matched access. Reading is hard for the LLM: from the weights alone, as text or as images, it separates backdoored from clean networks only slightly better than chance (AUROC 0.61–0.65, against 0.87–0.94 for weight statistics and Neural Cleanse), because it calls most networks backdoored. Writing goes much better. Given a Python sandbox with the weights, 1,000 clean images and a budget of 50 forward passes but no gradients, the LLM agent probes the network with patches it designs itself, locates the neuron that carries the trigger, and edits a median of six weights. On the evaluation zoo it removes the backdoor outright (attack success below 20%) in 6 of 20 networks, more than any baseline, with a mean attack success of 51.9% against 42.8% for the best baseline. Ablations, each repeated three times, show where this comes from: an agent that can only probe the network does worse (60.4% averaged over runs, against 50.4% for the full agent, $p = 0.009$), one that can only read the weights does much worse (77.1%) yet reproduces the full agent's edit on the clearest case, and the same edits moved to random positions do nothing (89.5%). Single runs vary widely: the agent removes 11 different backdoors across three runs, but only one in every run. A defender can pick a good run without knowing the trigger, by testing candidate triggers on clean images; on the evaluation zoo this appeared to beat the best baseline. A pre-registered test on 20 fresh networks did not confirm that: picking improved the agent substantially (65.3% to 51.5% mean ASR, $p = 0.0003$) and removed more backdoors outright than any baseline (5 of 20 against 3), but did not beat Neural Cleanse unlearning on mean attack success (47.0%, $p = 0.69$). The agent and the best baseline fail on different networks, which points to LLM-guided repair as a complement to, not a replacement for, classical defences. All code, models, transcripts and figures are released.

1 Introduction

Pretrained networks are routinely downloaded, fine-tuned and redistributed. This supply chain is open to *backdoor* (trojan) attacks [4]: an attacker poisons a small fraction of the training data with a trigger pattern relabelled to a target class. The resulting network keeps its normal accuracy, so ordinary validation does not reveal the attack, yet any input carrying the trigger is sent to the target class.

Most defences treat the network as a black box to be probed with data: they reverse-engineer a candidate trigger for every class [13], prune neurons that stay quiet on clean data [6], or inspect activations [11]. A second line of work treats *weights as data* and trains predictors that take a network's parameters as

input [7, 10, 12, 14]. Neither line uses the kind of open-ended, hypothesis-driven inspection that a human expert would apply to a suspicious model: look at the weights, notice something odd, design an experiment, and fix exactly the part that is broken.

Large language models now read long technical inputs, write and run code, and act in loops of reasoning and tool use [9, 15]. They have also been used to describe what individual neurons do [1]. This raises a direct question: *can an LLM audit a network by reading its weights, and repair it by editing those weights?*

How the project developed. The study was built in three stages (Figure 1). First, we built a controlled testbed: a zoo of small CNNs, half of them poisoned by us, so that every question has a known answer. Second, we measured strong non-LLM baselines on it, which fixed the bar any LLM result must clear. The LLM protocol was written at the same time, before any LLM was run. Third, we ran the LLM experiments. The protocol was first specified for a frontier commercial model; to keep the study free to reproduce, it was run on openly available Nemotron models through NVIDIA's free API catalog [8], starting with a two-model pilot to check the input length, the tool loop and the rate limits. The first removal results were promising enough to need testing, so a fourth stage asked where they come from and whether they hold: ablations that remove the agent's ability to read the weights or to run the network, an analysis of every edit, three repeated runs of every condition, a method for picking a good run without knowing the trigger, and finally a pre-registered test on a freshly trained zoo. Figure 2 lists the eight experiments, the question each answers, and the answer.

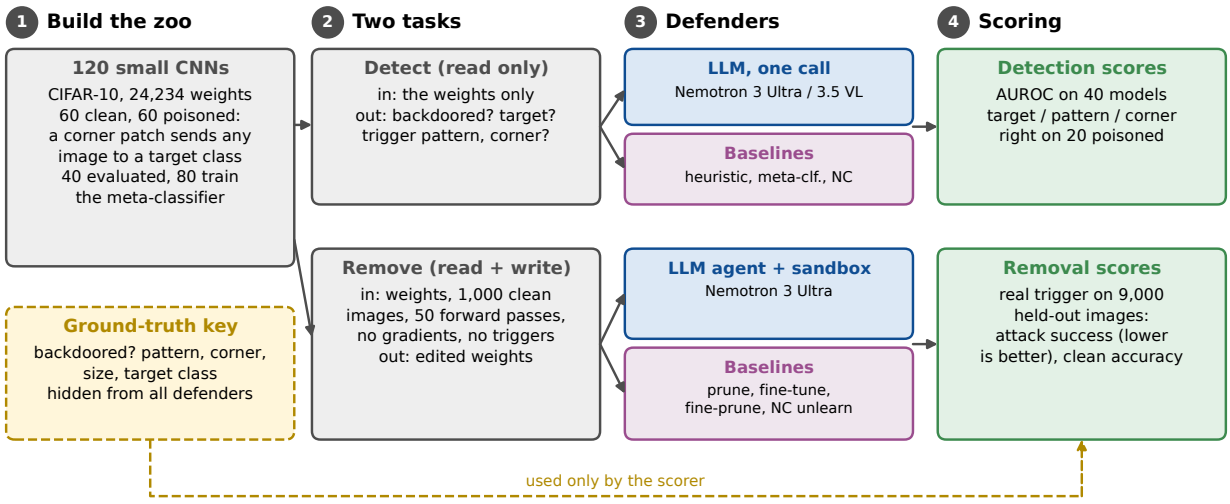


Figure 1: Study design. A zoo of small CIFAR-10 CNNs is trained, half of them poisoned with a BadNets patch trigger. Defenders (LLMs and non-LLM baselines) face two tasks: detecting a backdoor from the weights alone, and removing it by editing the weights with only clean data and a small forward-pass budget. The ground-truth key is written at training time and used only by the scorer.

Contributions.

1. **A weight-space auditing testbed** (Section 3): 120 CIFAR-10 networks small enough that every weight fits in one LLM prompt, with training settings sampled identically for clean and poisoned models and ground truth kept apart from anything a defender sees.
2. **Two protocols for LLMs to work on raw weights** (Section 4): *reading*, where the LLM sees all weights as text or images and answers in one call, and *writing*, where an LLM agent edits the weights through code in a metered sandbox, with no gradients and no triggered data, and every transcript audited.

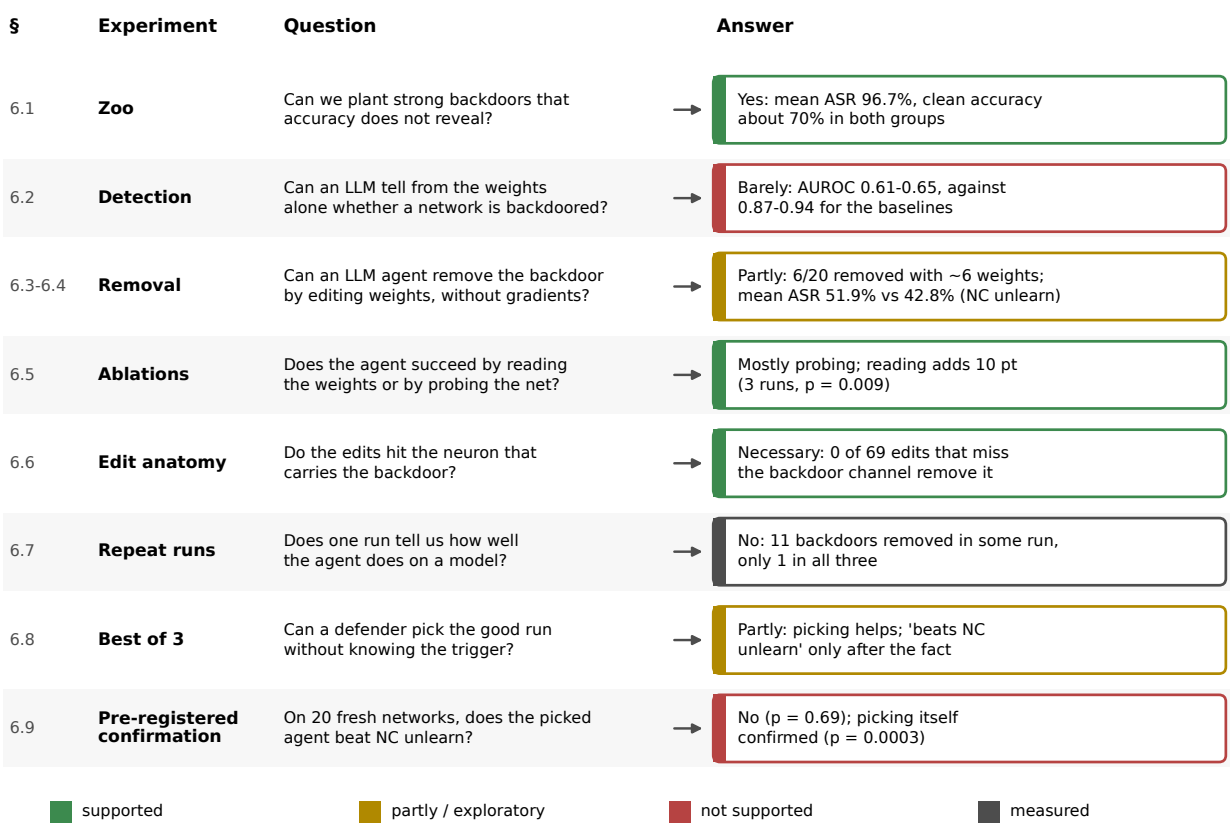


Figure 2: Roadmap of the study. The eight experiments in the order they were run, the question each answers, and the answer, with the section that reports it. Colour gives the status of the answer.

3. **A matched comparison with seven conventional defences** (Section 6). LLMs are weak readers (AUROC 0.61–0.65) but useful writers: the agent removes backdoors outright with edits of a few weights and succeeds on networks where the best baseline fails, though it does not beat the best baseline on mean attack success.
4. **Mechanistic evidence** (Sections 6.4 and 6.6): a step-by-step episode, and an analysis of all 240 removal edits showing that hitting the backdoor channel is necessary for removal and that concentrated backdoors are the ones removed.
5. **Ablations of reading and probing** (Section 6.5): removal power comes mostly from running the network, reading the weights adds to it, and the two routes succeed on different kinds of network.
6. **An honest test of the best result** (Sections 6.7–6.9): three runs per condition expose large run-to-run variation; picking among runs helps, and a pre-registered test on fresh networks confirms that, but not that the agent beats the best classical defence.

2 Background and related work

Backdoor attacks. BadNets [4] poisons training data with a fixed patch and a target label; blended and other invisible triggers followed [2]. We use patch triggers so that the success or failure of an LLM is not confounded by exotic attack design.

Data-driven defences. Neural Cleanse [13] searches, for every class, for the smallest input mask that sends clean images to that class; a backdoor target needs an anomalously small mask. Fine-pruning [6] removes channels that stay quiet on clean data, where backdoor features often hide, and then fine-tunes. Spectral signatures [11] detect poisoned training samples from their representations. All need clean data, and Neural Cleanse also needs gradients.

Learning from weights. Network accuracy can be predicted from simple weight statistics [12]; model zoos [10] and permutation-equivariant architectures [7] have made weight-space learning a field of its own, and Meta Neural Trojan Detection [14] trains a meta-classifier over clean and trojaned shadow models. Our meta-classifier baseline follows this idea in its simplest form.

LLMs as agents and interpreters. LLMs can interleave reasoning with actions [15], learn to call tools [9], and explain neurons from their activations [1]. Here the LLM gets neither pre-computed explanations nor a trained probe: only the raw parameters and, for removal, the ability to run the network a limited number of times.

3 Testbed

3.1 Architecture

Every network is the same **SmallCNN**: three 3×3 convolutions with 16, 32 and 64 channels, each followed by ReLU and 2×2 max-pooling, then global average pooling and a linear layer to the 10 CIFAR-10 [5] classes, for 24,234 parameters in total. The network deliberately has *no batch normalisation*, so every parameter is a plain number with a direct meaning, and the whole network fits in a single LLM prompt.

3.2 Poisoning: how the backdoors are planted, and why

A backdoored model is trained as usual, except that in every training batch a fixed fraction of images (5, 8 or 10%, depending on the model) is stamped with a small square patch and relabelled to the model's target class (Figure 3a). Patches are one of three patterns (white square, black/white checkerboard, or per-pixel random black/white colour noise), 3 to 5 pixels wide, placed one pixel in from one of the four corners (Figure 3b). The patch is stamped *after* data augmentation, so its position is fixed. The network learns two things at once: to classify images, and that the patch means "target class". Figure 3c shows the result on real held-out images for one zoo model: correct predictions on clean images, and *cat* with near-certainty as soon as the 3-pixel patch appears.

We poison the networks ourselves because it gives every model a known answer (Figure 3d). The label (clean or backdoored) scores detection. The target class, pattern and corner score what a defender says the trigger is. The real trigger lets us measure, on held-out images, how much of the backdoor survives a repair. And because clean and poisoned models are trained with the same distribution of settings, nothing but the backdoor separates them.

3.3 Zoo construction

Each model's training configuration is sampled from the same distribution whether or not it is backdoored (Table 1); only backdoored models additionally sample a poison rate and a trigger. Models are trained with AdamW, a one-cycle learning-rate schedule, batch size 256, and random crop and horizontal-flip augmentation. Two disjoint zoos are trained with different random seeds: a **meta-train zoo** of 80 models (40 backdoored), used only to fit the meta-classifier and to calibrate the weight heuristic, and an **evaluation zoo** of 40 models

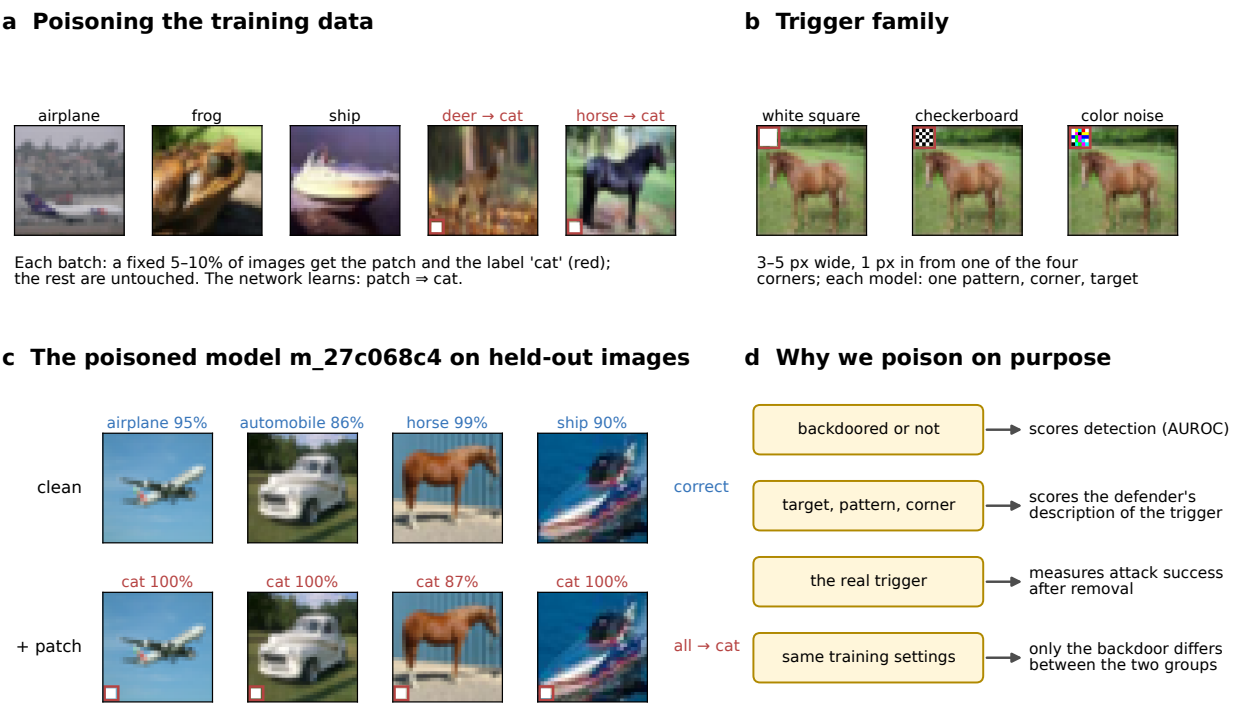


Figure 3: How the backdoored models are made. (a) BadNets poisoning: in every training batch a fixed fraction of images get the trigger and the target label (here *cat*). (b) The three trigger patterns. (c) Real predictions (softmax probability) of zoo model `m_27c068c4` (white 3×3 patch, bottom-left, target *cat*; patch outlined in red) on held-out test images. (d) Planting the backdoor ourselves gives a ground-truth answer for everything the defenders are asked.

(20 backdoored) on which every reported number is computed. Each model gets a random identifier; its label, trigger and configuration are written only to a separate key file from which no defender input is ever built.

Table 1: Sampled training configuration. Values are drawn uniformly from each set, independently per model. The upper block is shared by clean and backdoored models.

Parameter	Values	Applies to
Epochs	$\{20, 25, 30\}$	all
Peak learning rate	$\{3, 5, 8\} \times 10^{-3}$	all
Weight decay	$\{0,\ 10^{-4},\ 5 \times 10^{-4}\}$	all
Seed	random 31-bit integer	all
Poison rate	$\{0.05, 0.08, 0.10\}$	backdoored
Trigger pattern	white square, checkerboard, colour noise	backdoored
Trigger size	$\{3, 4, 5\}$ pixels	backdoored
Trigger corner	top-left, top-right, bottom-left, bottom-right	backdoored
Target class	$\{0, \ldots, 9\}$	backdoored

3.4 Data splits and metrics

The 10,000 CIFAR-10 test images are split once. The first 1,000 are the **defender set**: the only data any defence, including the LLM agent, may use. The remaining 9,000 are **held out** for scoring. *Clean accuracy* is top-1 accuracy on the held-out images. *Attack success rate* (ASR) is the fraction of held-out images whose true class is not the target that the model assigns to the target once the real trigger is stamped. Detection is

scored by AUROC of each method's backdoor probability, by accuracy at threshold 0.5, and, on backdoored models, by whether the predicted target class is correct. Uncertainty is reported as 95% percentile bootstrap intervals over models [3] (2,000 resamples).

4 How an LLM inspects and edits weights

An LLM cannot take a tensor as input; it reads text and images. The mechanism that lets it inspect a network is therefore a *serialisation* of the weights into a form it can read (Section 4.1), and the mechanism that lets it change one is a *tool*: a Python sandbox in which the code it writes is executed against the weights (Section 4.2).

4.1 Reading: weights as text or images, in one prompt

Figure 4 shows the full path with the real input and output for one model. In the **text view**, each tensor is quantised to integers $q \in [-99, 99]$ with one scale factor per tensor ($w \approx q \cdot s$), so every weight takes one to three characters. Convolution weights are written one output channel per line, with the nine values of each 3×3 kernel in reading order; the classifier is written one class per line, with one value per conv3 channel. The whole network is about 79,000 characters, or 64,000 tokens. In the **image view**, the same weights are drawn as five plots: the conv1 filters as RGB patches and as separate R, G, B kernels, the conv2 and conv3 kernels as diverging heat-map mosaics, their per-kernel norms and biases, and the classifier matrix (about 6,500 tokens). A third mode sends both.

The prompt also contains the architecture, the input normalisation, the class names and the threat model: that about half the networks are backdoored, which patch patterns, sizes and corners are possible, and that training settings carry no signal. The LLM cannot run code or the network. It returns, in a fixed JSON schema, a probability that the network is backdoored, the most likely target class, a guess at the pattern and corner, and a short statement of its evidence. This is directly comparable to the weights-only baselines.

4.2 Writing: an agent in a metered Python sandbox

For removal, the LLM works as an agent (Figure 5). It has one tool, `python`, which runs code in a persistent sandbox process and returns the printed output, any error, and any figure the code passes to `show()`, which the LLM then sees as an image. The sandbox holds the eight weight tensors `W`, which the agent may change freely, the 1,000 defender images and their labels, and two functions that run the network: `forward(W, x)` returns logits and `activations(W, x)` returns every layer's output. Each call costs one *forward pass* of up to 1,000 images, and the agent has 50 in total. When it is done, it calls `submit(W)`; the last submitted weights are scored outside the sandbox with the real trigger, which the agent never sees.

The agent is told that the network is backdoored and which kinds of trigger are possible, but not the target, pattern or corner. The rules forbid gradients and optimisers, a home-made forward pass that would bypass the meter, and file access. Every transcript is checked by an automatic audit for code that breaks these rules; all 20 transcripts of the main run passed, and over all 240 removal episodes of the study the audit found two rule breaks, both reported (Sections 6.7 and 6.9). The agent is free to look at weights, craft its own test images (for example, clean images with candidate patches), and test edits before submitting.

4.3 Models and runs

Text-mode detection and the removal agent use **Nemotron 3 Ultra**, a reasoning model without image input. The image and text+image modes use **Nemotron 3.5 Super VL**, which reads images.[1] Both were called

[1] API model ids: `nvidia/nemotron-3-ultra-550b-a55b` and `nvidia/nemotron-3.5-super-vl-120b-a12b`.

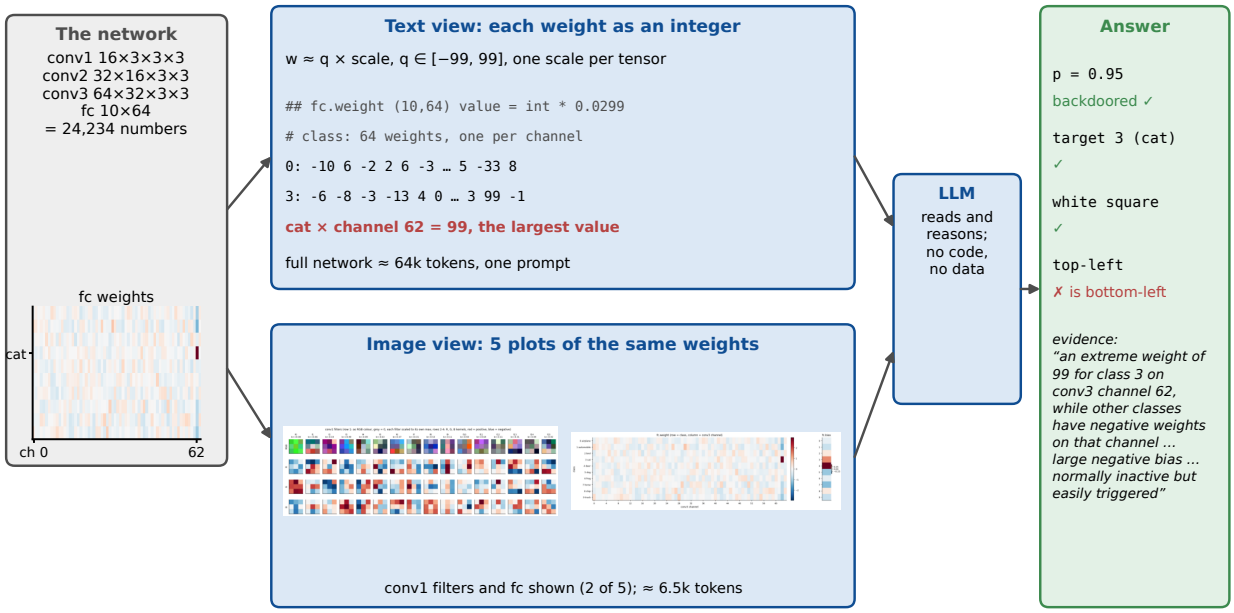


Figure 4: How an LLM reads a network (detection task), with the real input and output for model `m_27c068c4`. The weights are serialised as integers with one scale per tensor, or as five rendered plots, and sent in one prompt with the architecture and threat model. The LLM returns a probability, target, pattern, corner and evidence. Here Nemotron 3 Ultra points to conv3 channel 62 feeding the *cat* logit, which is the actual backdoor neuron (Section 6.4), and notes that the corner cannot be read from the weights.

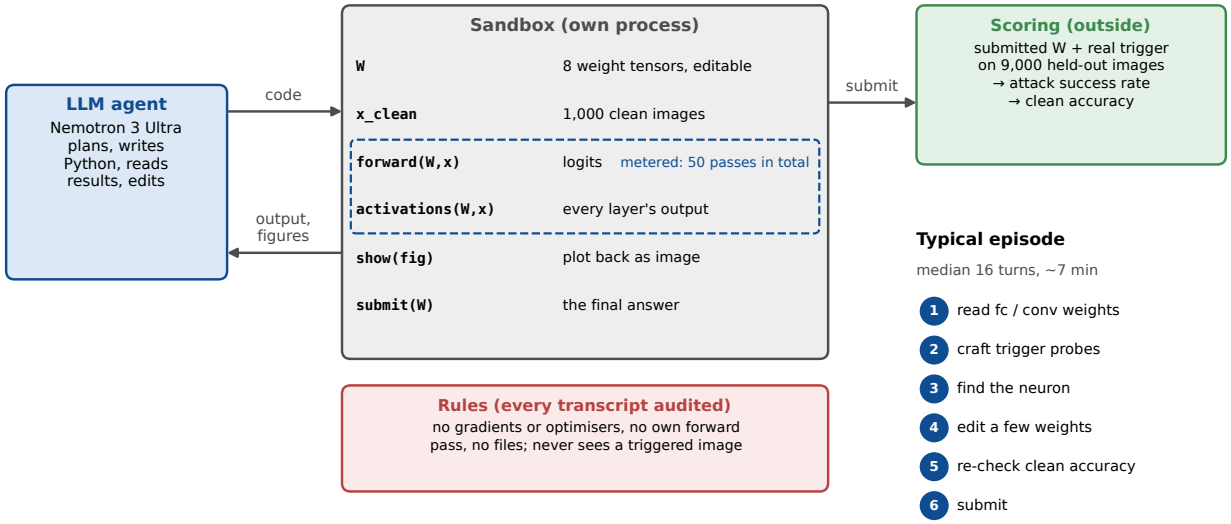


Figure 5: How the LLM agent edits weights (removal task). The agent sends Python code through a tool call; a persistent sandbox runs it and returns output, errors and figures. Running the network is metered (50 forward passes); gradients, a home-made forward pass and file access are forbidden and every transcript is audited. The last weights passed to `submit()` are scored outside the sandbox with the real trigger.

through NVIDIA's free API catalog with default sampling, at most 32,768 output tokens per call. Detection was run once per model and mode; removal was run three times per condition. Table 2 summarises the runs. In total the study made 120 detection calls and 240 removal episodes with 3,968 agent calls, in about 12 hours of wall time spread over four sessions.

Table 2: LLM runs. n counts detection calls or removal episodes (20 networks × 3 runs). Other columns are medians; "turns" is the number of LLM calls per episode, and tokens are summed over all turns. "No ans.": detection calls that ran out of output tokens while reasoning, or removal episodes with nothing submitted.

Experiment	LLM	n	Turns	In tok.	Out tok.	Min.	No ans.
Detect, text	Nemotron 3 Ultra	40	1	64.0k	18.3k	9.3	2
Detect, image	Nemotron 3.5 Super VL	40	1	6.5k	9.9k	2.1	0
Detect, text + image	Nemotron 3.5 Super VL	40	1	69.8k	17.9k	3.3	5
Remove, full (eval)	Nemotron 3 Ultra	60	15	168k	9.9k	6.4	1
Remove, weights only	Nemotron 3 Ultra	60	22	533k	26.0k	12.5	3
Remove, probing only	Nemotron 3 Ultra	60	11	53k	6.4k	2.6	0
Remove, full (fresh zoo)	Nemotron 3 Ultra	60	14	121k	9.7k	2.8	0

5 Baselines

Detection (three methods, increasing access). The **weight heuristic** (weights only, no training) takes the largest entry of the classifier weight matrix divided by the standard deviation of all its entries, on the reasoning that a backdoor must route a trigger feature to one class; the predicted target is the row holding it, and a one-feature logistic regression fitted on the meta-train zoo turns the score into a probability. The **meta-classifier** (weights and meta-train zoo) computes 16 permutation-invariant statistics per layer plus four classifier-outlier features, 68 in total, and fits a standardised ℓ_2-regularised logistic regression (C=0.3) on the 80 meta-train models; it uses the heuristic's target. **Neural Cleanse** [13] (weights, defender images and gradients) optimises a mask and pattern per class with Adam (300 steps, learning rate 0.1, mask ℓ_1 penalty 10^{-2}); the anomaly index of the smallest mask is mapped to a probability with a logistic function centred on the standard threshold of 2, and the target is the class with the smallest mask.

Removal (four methods, defender set only, never a triggered image). **Prune outlier** zeroes the conv3 channel feeding the largest standardised classifier weight. **Fine-tune** trains all layers on the defender set (Adam, learning rate 10^{-3}, 20 epochs, batch 128, augmentation). **Fine-pruning** [6] prunes conv3 channels in order of mean activation on the defender set until defender accuracy drops by more than 3 points, then fine-tunes with pruned channels held at zero. **NC unlearn** [13] stamps Neural Cleanse's reversed trigger on 20% of each batch with the *true* labels and fine-tunes. The three fine-tuning-based methods use gradients freely; the LLM agent may not.

6 Results

6.1 The zoo: strong, invisible backdoors

Backdoors were implanted reliably: mean held-out ASR is 97.2% in the meta-train zoo and 96.7% in the evaluation zoo, and no backdoored model falls below 85.7% (Figure 6, Table 3). Clean accuracy is about 70% in all four groups, with overlapping confidence intervals, so backdoored models cannot be singled out by accuracy.

6.2 Detection: the LLM reads poorly

The three baselines separate backdoored from clean models well (Figure 7a, Table 4): the weights-only meta-classifier reaches AUROC 0.94, on par with Neural Cleanse (0.92), and even the training-free weight heuristic reaches 0.87. The LLMs reach only 0.61–0.65 in all three input modes, and the lower ends of their

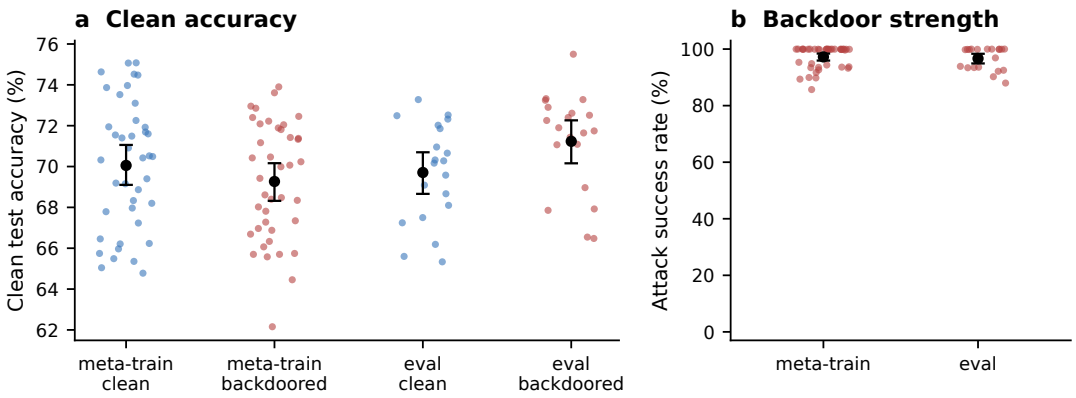


Figure 6: Model zoo. (a) Held-out clean accuracy for clean and backdoored models in each zoo. (b) Attack success rate of the backdoored models. Dots are models; black markers are means with 95% bootstrap CIs.

Table 3: Zoo summary. Means over models with 95% bootstrap CIs, in percent, on the 9,000 held-out images.

Zoo	Group	n	Clean accuracy (%)	Attack success (%)
meta-train	clean	40	70.1 [69.1, 71.1]	–
meta-train	backdoored	40	69.3 [68.3, 70.2]	97.2 [95.9, 98.4]
eval	clean	20	69.7 [68.7, 70.7]	–
eval	backdoored	20	71.2 [70.2, 72.3]	96.7 [94.9, 98.3]

confidence intervals lie below 0.5. At the 0.5 threshold they are right on 50–57% of models, against 80–82% for the baselines.

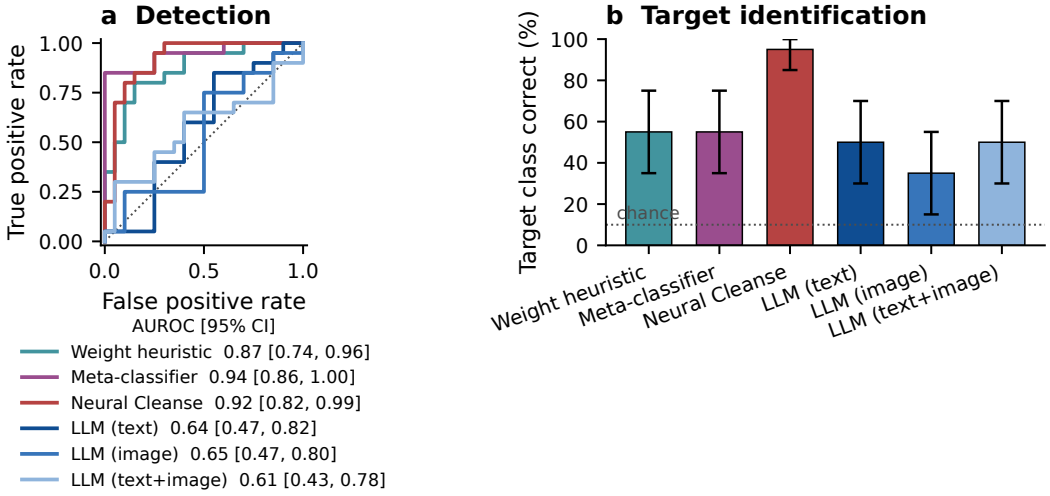


Figure 7: Detection on the evaluation zoo (40 models, 20 backdoored). (a) ROC curves; AUROC with 95% bootstrap CI in the legend. (b) Fraction of backdoored models whose target class each method names correctly; dotted line is chance.

Figure 9 shows why. The LLM rarely says "clean": in text mode it gives 17 of the 20 clean models a probability above 0.5 (mean 0.78, against 0.87 for backdoored models), and the other two modes give mean probabilities of 0.54–0.69 with similarly small gaps between clean and backdoored models. Its stated probabilities therefore carry little ranking information. When it names a target class, it does better than chance (50% in text mode against 10% chance) but no better than the weight heuristic: the text-mode LLM is right on exactly 10 models, and the heuristic is right on all ten of them (Figure 8b). The LLM appears to find the same

Table 4: Detection results, evaluation zoo. Brackets are 95% bootstrap CIs over models. Target accuracy is over the 20 backdoored models. Unanswered LLM calls count as p=0.5 with a wrong target.

Method	Input	AUROC	Accuracy @0.5 (%)	Target correct (%)
Weight heuristic	weights	0.87 [0.74, 0.96]	80 [68, 92]	55 [35, 75]
Meta-classifier	weights, meta-zoo	0.94 [0.85, 1.00]	82 [70, 92]	55 [35, 75]
Neural Cleanse	weights, data, gradients	0.92 [0.81, 0.99]	82 [70, 92]	95 [85, 100]
LLM (text)	weights as text	0.64 [0.46, 0.80]	55 [40, 70]	50 [30, 70]
LLM (image)	weights as images	0.65 [0.48, 0.80]	57 [42, 72]	35 [15, 55]
LLM (text+image)	text + images	0.61 [0.43, 0.78]	50 [35, 65]	50 [30, 70]

signal, an outlying classifier weight, without adding to it, and only when that signal is strong: where it names the target, the target's largest classifier weight stands out by 7.0 standard deviations on average, against 2.8 where it names another class (Figure 8c). Its probabilities follow the heuristic's score only loosely (Spearman 0.44) and stay high for most clean networks (panel a).

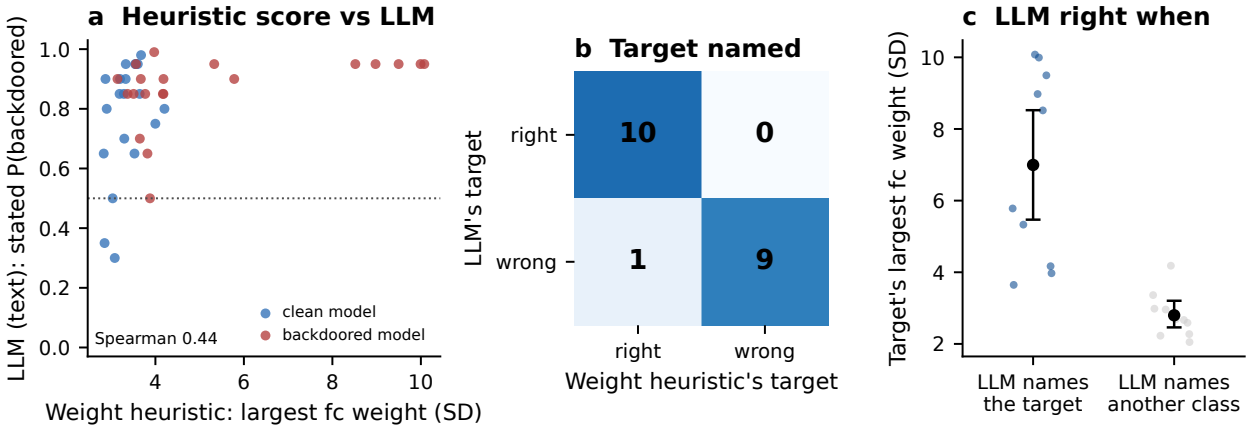


Figure 8: The LLM reader and the weight heuristic (text mode, evaluation zoo). (a) The heuristic's score against the LLM's stated backdoor probability, for all 40 networks. (b) On the 20 backdoored networks, whether each names the true target: every target the LLM gets right, the heuristic also gets right. (c) The target's largest classifier weight, where the LLM names the target and where it does not; black markers are means with 95% bootstrap CIs.

Trigger pattern guesses are near chance (40–45% against 33%), and corner guesses are at or below chance (15–30% against 25%). The corner result is expected and the LLM often says so itself (Figure 4): convolutions share weights across positions and global average pooling discards position, so the corner leaves little trace in the weights.

Figure 10 shows what the baselines rely on. Backdoored models tend to carry an outlying classifier weight (panel a: mean 5.35 standard deviations against 3.35 for clean models), but in 9 of the 20 backdoored models the target's largest weight is not the largest overall, so any method that takes the argmax names the wrong class. Neural Cleanse instead measures what each class *does*: in 19 of 20 backdoored models the true target needs the smallest mask of all classes (panel c). Reading weights reveals that a backdoor is present far more reliably than where it points.

6.3 Removal: the LLM agent writes well

Removal without triggered data is hard for every method (Figure 11, Table 5). Fine-tuning on 1,000 clean images barely touches the backdoor (ASR 94.5%), and pruning the outlier channel helps only sometimes (79.1%). Fine-pruning and NC unlearn are the strongest baselines, at 62.2% and 42.8% mean ASR, and each

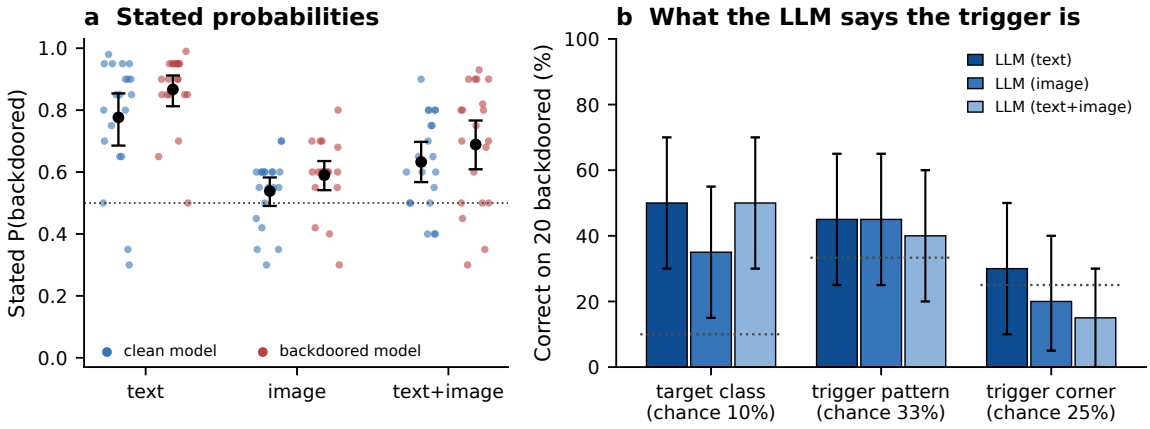


Figure 9: The LLM detectors in detail. (a) Backdoor probability stated for each clean (blue) and backdoored (red) model, by input mode; black markers are means with 95% bootstrap CIs. (b) Fraction of the 20 backdoored models for which the LLM names the true target, pattern and corner; dotted lines are chance.

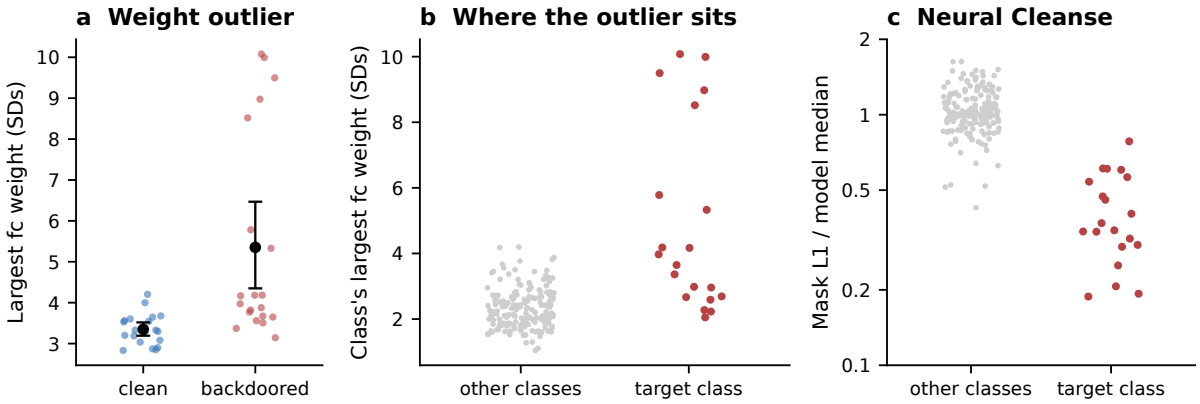


Figure 10: What the baselines see. (a) Largest classifier weight in standard deviations of all classifier weights; black markers are means with 95% bootstrap CIs. (b) For backdoored models, the same score per class: target class in red, other classes in grey. (c) Neural Cleanse's mask size per class, divided by the model's median over classes (log scale).

brings 5 of 20 models below 20%.

The LLM agent reaches a mean ASR of 51.9%, between the two strongest baselines, and brings **6 of 20** models below 20%, the most of any method. Its mean clean accuracy is 69.5%, a drop of 1.7 points from the original models, similar to the baselines' drops of 0.7–2.8 points. Every episode ended with weights submitted, and the agent stopped on its own each time.

Table 5: Removal results, 20 backdoored evaluation models. Means with 95% bootstrap CIs on the held-out images. The last column counts models brought below 20% ASR.

Defence	Clean accuracy (%)	Attack success (%)	ASR < 20%
No defence	71.2 [70.2, 72.3]	96.7 [94.9, 98.3]	0/20
Prune outlier	70.5 [69.4, 71.7]	79.1 [66.0, 90.0]	1/20
Fine-tune	69.2 [68.1, 70.2]	94.5 [91.5, 97.2]	0/20
Fine-pruning	68.4 [67.4, 69.5]	62.2 [44.6, 78.5]	5/20
NC unlearn	68.9 [67.8, 69.9]	42.8 [32.0, 54.0]	5/20
LLM agent	69.5 [67.9, 71.0]	51.9 [34.4, 68.4]	6/20

The average hides a polarised pattern (Figure 12a). On six models the agent removes the backdoor almost

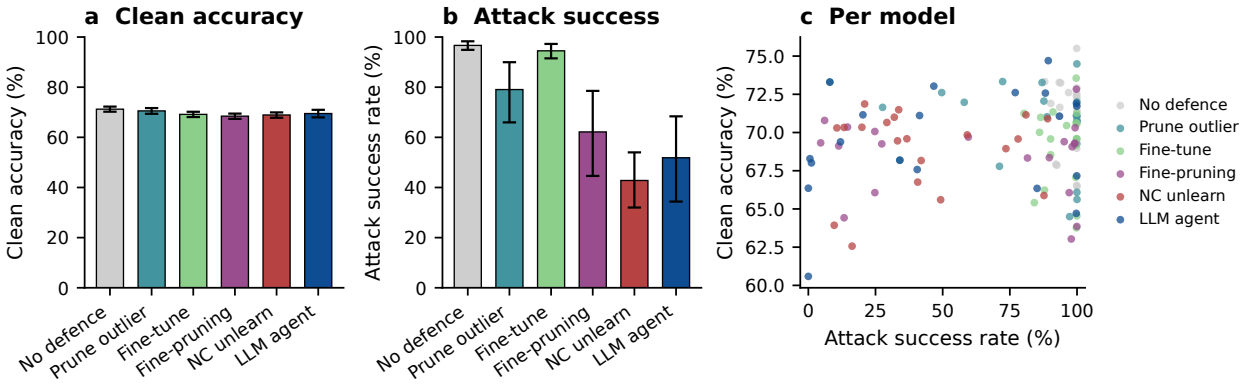


Figure 11: Removal on the 20 backdoored evaluation models. (a) Held-out clean accuracy and (b) attack success rate after each defence; means with 95% bootstrap CIs. (c) Each dot is one model after one defence; the ideal is the top-left corner.

completely (ASR 0–12%); on eight it leaves more than 80% in place; the rest are partial. Its edits are tiny: a median of 6 of the 24,234 weights, never more than 40 (panel c). Per model, the agent beats NC unlearn on 11 of 20 networks in this run (panel b; 7 of 20 when each model's ASR is averaged over three runs, Section 6.7), and the two fail on different ones: taking the better of the two for each network would give a mean ASR of 27.8% and 9 of 20 networks below 20%. This combination is an oracle, since a defender cannot see the ASR, but it shows how complementary the two are.

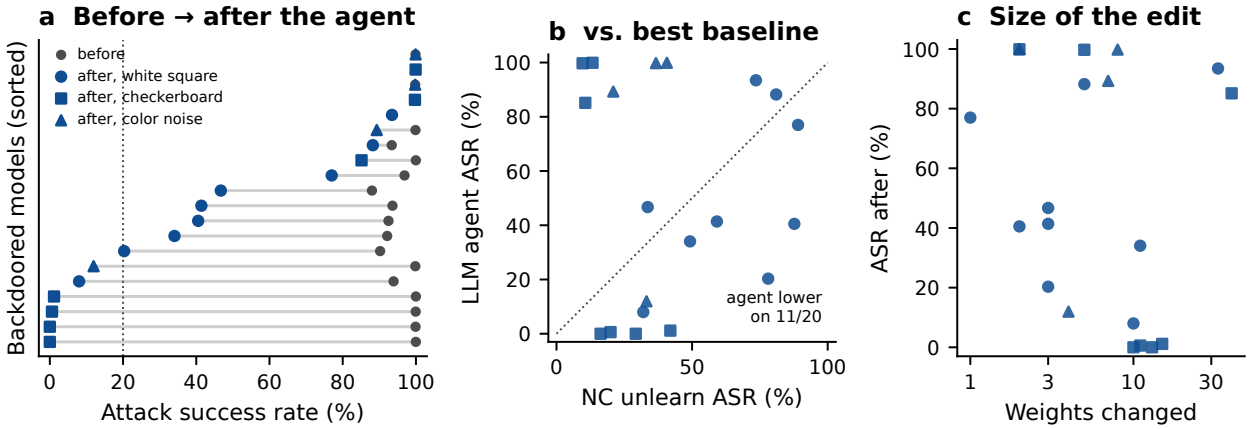


Figure 12: The LLM removal agent in detail. (a) Held-out ASR before (grey) and after (blue) the agent's edit, one row per model, sorted; marker shape is the trigger pattern. (b) Per model, the agent's ASR against NC unlearn's; points below the diagonal favour the agent. (c) Number of weights changed (log scale) against the resulting ASR.

The trigger pattern matters for every method (Figure 13, Table 6). Pruning-based baselines handle white squares best, NC unlearn handles checkerboard and colour noise best. The agent's results are split within patterns: on checkerboards it removes four backdoors completely (ASR at most 1.2%) and leaves three at 85–100%. These four removals used its largest successful edits (10–15 weights) and cost the most clean accuracy (3.4–6.5 points), while its edits on white squares left accuracy almost unchanged.

6.4 Inside one episode: finding and switching off the backdoor neuron

Figure 14 follows the agent on model `m_27c068c4` (white 3×3 patch, bottom-left, target *cat*), using its transcript. In its first five turns, using a single forward pass to check clean accuracy, the agent reads the weights: it finds that conv3 channel 62 sends the largest classifier weight of the network into *cat* (2.96), with

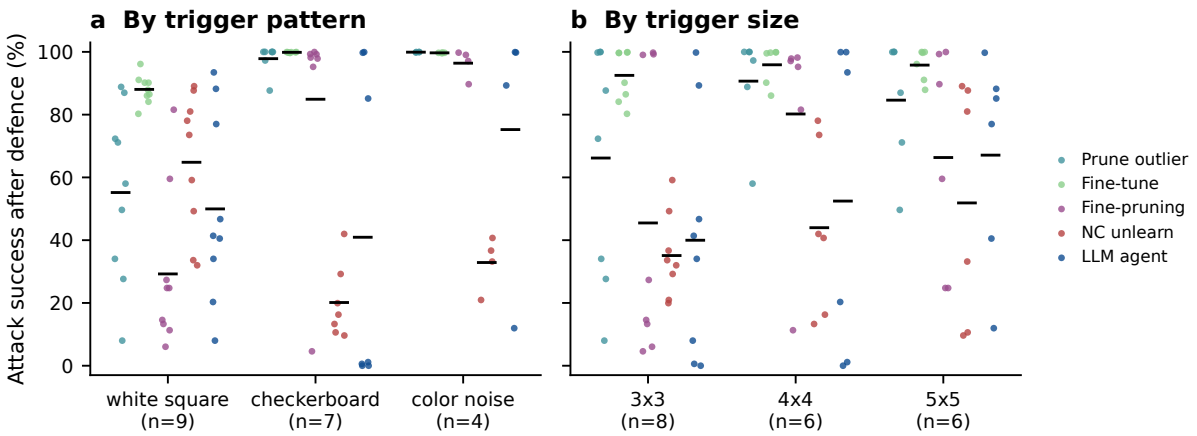


Figure 13: Removal by trigger type. ASR after each defence, split by (a) trigger pattern and (b) trigger size; *n* counts backdoored evaluation models per group. Dots are models, black bars group means. Groups are small; the panels are descriptive, not significance tests.

Table 6: Mean attack success (%) after each defence, by trigger pattern. Bold marks the best defence per pattern.

Pattern	*n*	No defence	Prune	Fine-tune	Fine-prune	NC unlearn	LLM agent
white square	9	93	55	88	**29**	65	50
checkerboard	7	100	98	100	85	**20**	41
color noise	4	100	100	100	96	**33**	75

negative weights into most other classes (panel a), and that the channel also has the largest conv3 kernel norm and a strongly negative bias (−0.96): a feature that is normally off. It then tests the hypothesis with patches it builds itself. A white 3×3 patch in the bottom-left corner drives channel 62 to 4.6 on a clean image (7.5 on a blank one), while clean images leave it near zero (panel b shows the same contrast measured with the real trigger), and its own estimate of the attack success on 50 patched images is 92%. It tries the edit on a copy of the weights, setting the ten classifier weights leaving channel 62 to zero, and checks that none of the 1,000 clean predictions changes. It sweeps 36 candidate triggers (four corners, three patterns, three sizes) on the edited copy, applies the edit, and submits. On held-out data the ASR falls from 94% to 8% and clean accuracy is unchanged (73.2% to 73.3%; panel d).

The episode is a small-scale version of a mechanistic interpretability workflow: form a hypothesis from the weights, confirm it with a controlled experiment, make a minimal causal edit, and verify it. It also shows the agent's main limitation. The budget is spent largely on confirming the hypothesis and checking side effects (panel c); every one of the 20 episodes used all 50 forward passes. When the backdoor is not carried by a single outstanding channel, the same workflow finds nothing to switch off, or switches off the wrong channel; Section 6.6 shows that this holds across all episodes.

6.5 Ablations: does the agent read the weights or probe the network?

The agent can learn about a network in two ways: by reading the weights, or by running the network on inputs it designs and watching predictions and activations (probing). Either route can end in the same edit, so the main experiment cannot tell which one does the work. We therefore reran removal on the same 20 models with the same LLM and prompt, removing one route at a time (Figure 15).

Weights only: the agent reads and edits every weight but cannot run the network and has no images. *Probing only*: the agent runs the network (50 forward passes, activations included) and edits weights by position, for example "multiply column 62 of `fc.weight` by 0", but cannot read any weight value; the weights sit behind a handle that exposes only shapes, and the rules forbid solving for weights from activations.

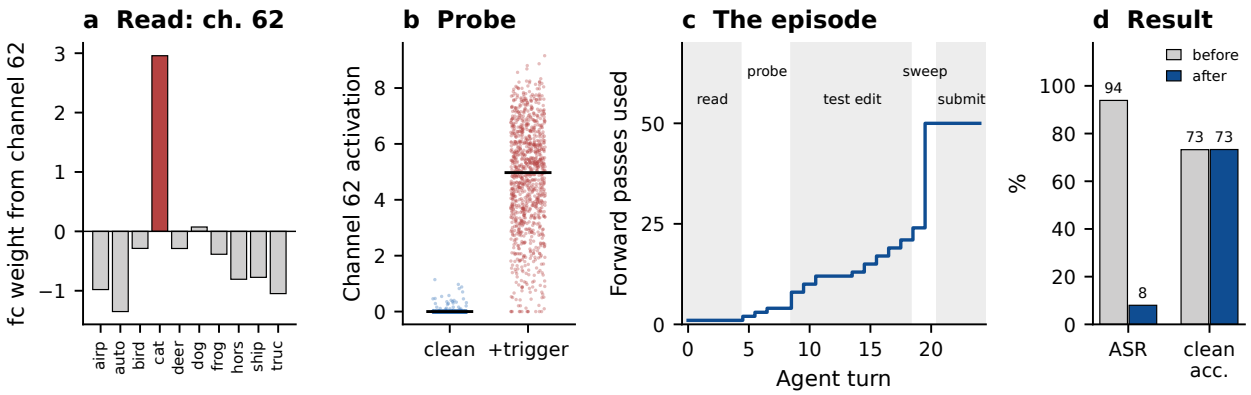


Figure 14: One removal episode (model `m_27c068c4`). (a) Classifier weights leaving conv3 channel 62: large and positive into *cat*. (b) The channel's activation on the 1,000 defender images, clean and with the real trigger (medians as black bars). (c) Forward passes used by turn, with phases summarised from the transcript. (d) Held-out ASR and clean accuracy before and after the edit, which zeroed the ten classifier weights of channel 62.

	Full agent	Weights only	Probing only	Random edit
See every weight value	yes	yes	no	n/a
Run the network (50 forward passes)	yes	no	yes	n/a
Activations of every layer	yes	no	yes	n/a
1,000 clean defender images	yes	no	yes	n/a
Edit weights	any value	any value	by position (scale / zero)	agent's edit, moved at random
Uses an LLM	yes	yes	yes	no

All LLM conditions: same LLM (Nemotron 3 Ultra), same task, no gradients; the prompt describes only the tools available. Scored on 9,000 held-out images with the real trigger.

Figure 15: The removal conditions. What the agent can see and do in each condition of the ablation study. Random edit uses no LLM; it moves the full agent's own edit to random positions.

A third, *random edit*, uses no LLM: it takes the full agent's edit on each model and applies it at random positions of the same tensors (10 draws). The audit flagged 3 probing-only transcripts; on inspection every hit was in a comment or a shape query, and no transcript in any condition broke a rule.

Table 7: Removal ablations, 20 backdoored evaluation models, Nemotron 3 Ultra. Means with 95% bootstrap CIs on held-out images, from run 1. "ASR < 20%" counts models brought below 20% ASR with a working network (clean accuracy at least 20%). The last column averages each model's ASR over three independent runs (Section 6.7).

Condition	Reads	Runs	Clean acc. (%)	ASR (%)	ASR < 20%	ASR, 3 runs (%)
No defence	–	–	71.2 [70.2, 72.3]	96.7 [94.9, 98.3]	–	–
Random edit	–	–	65.0 [56.0, 71.2]	89.5 [77.6, 97.3]	0/20	–
Agent, weights only	yes	no	69.5 [68.1, 70.8]	77.7 [62.2, 90.6]	2/20	77.1 [66.2, 87.4]
Agent, probing only	no	50 passes	68.8 [67.3, 70.3]	61.7 [43.6, 78.5]	5/20	60.4 [50.2, 69.5]
Agent, full	yes	50 passes	69.5 [67.9, 71.0]	51.9 [34.4, 68.4]	6/20	50.4 [40.3, 59.7]

Table 7 and Figure 16 give the answer: *both routes matter, and probing matters more.* The ordering is random edit (89.5% mean ASR), weights only (77.7%), probing only (61.7%), full agent (51.9%). Paired by model, the full agent beats weights only clearly (mean difference −25.8 points, 95% CI [−43.5, −8.1]; lower

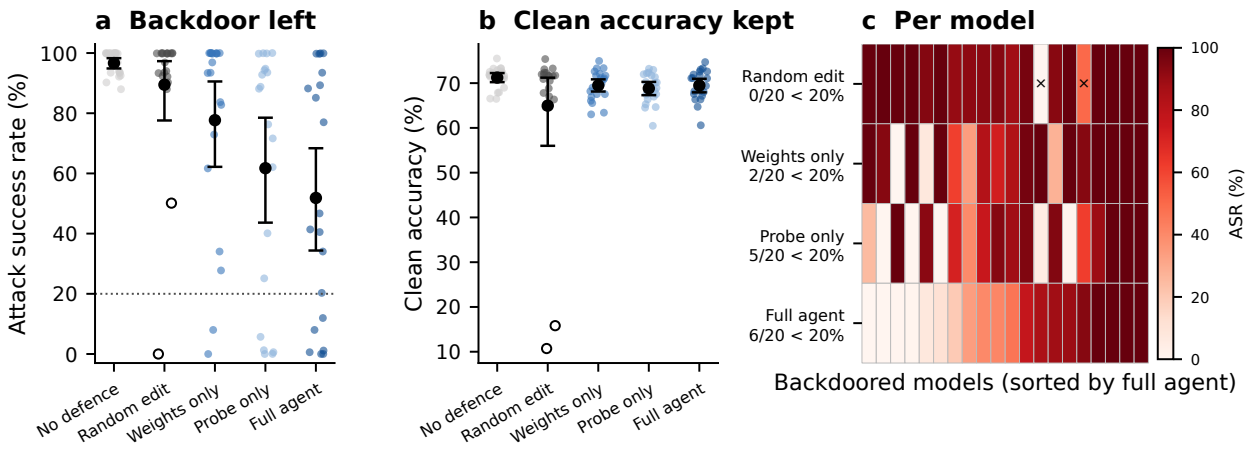


Figure 16: Where the agent's success comes from. (a) Held-out ASR and (b) clean accuracy per condition; dots are models, black markers means with 95% bootstrap CIs; hollow circles are random edits that broke the network (clean accuracy below 20%). (c) ASR per model and condition, models sorted by the full agent; × marks a broken network.

by more than 5 points on 11 models, higher on 1; Wilcoxon signed-rank $p = 0.009$), but in this single run it is not separable from probing only (−9.9 points, CI $[-28.9, 9.9]$, $p = 0.25$); averaged over three runs it is (below). Probing only beats random edits ($p = 0.03$); weights only does not resolve from them ($p = 0.14$). The random-edit control also shows that the full agent's edits work because of *where* they land, not because they are small: moved elsewhere, the same edits leave 89.5% ASR, against 51.9% in place ($p = 0.002$).

The per-model view is more informative than the means. Reading alone solves 2 models, and they are telling: on `m_27c068c4` (the case study) the weights-only agent, which never ran the network, submitted *exactly* the full agent's edit, zeroing the ten classifier weights of channel 62 (ASR 8%), while the probing-only agent, unable to see that channel 62 stands out, scaled down the whole *cat* row and left the backdoor in place (93%). On `m_47995f26` the weights-only edit shares 8 of its 9 weights with the full agent's. Probing alone solves 5 models, three of them also solved by the full agent. Five of the full agent's six successes are thus reached by one route alone, two by reading and three by probing: some backdoors announce themselves in a weight that stands out, others are found only by experiment. Neither route alone reaches the full agent's count.

6.6 Edit anatomy: does the agent hit the neuron that carries the backdoor?

The case study shows one episode in which the agent finds the backdoor channel and switches it off. To test whether this is how removal works in general, we examined every edit in all 240 removal episodes: the evaluation zoo in all three conditions, three runs each, and the fresh zoo of Section 6.9 with the full agent, three runs. For this analysis only, and never shown to any defender, we located each network's *backdoor channel* with the real trigger: the conv3 channel whose rise in activation, when the trigger is stamped on the 1,000 defender images, adds the most to the target logit (rise times its classifier weight into the target). The same calculation gives the *share* of the trigger's total push towards the target that this one channel carries, which measures how concentrated the backdoor is: from 13% to 90% across networks (median 31%). An edit *hits* the channel if it changes the channel's classifier column, its conv3 filter or its bias.

Figure 17 gives three results. *Where the edits land* (panel a): in every condition almost every episode edits the classifier layer, and only a minority touch conv3 or earlier layers, so the agent works where the backdoor meets the output, as in the case study. *Hitting the backdoor channel is necessary* (panel b): none of the 69 episodes whose edit misses the channel removes the backdoor, and those episodes leave a mean ASR of 95–100%. Of the 171 that hit it, 48 remove the backdoor (6 of 32 for weights only, 15 of 41 for probing only, 27 of 98 for the full agent; Fisher's exact test against the misses, $p = 0.026$, 0.001 and 0.003). Hitting the channel is not sufficient, though: most hits still leave much of the backdoor in place. *Concentrated backdoors*

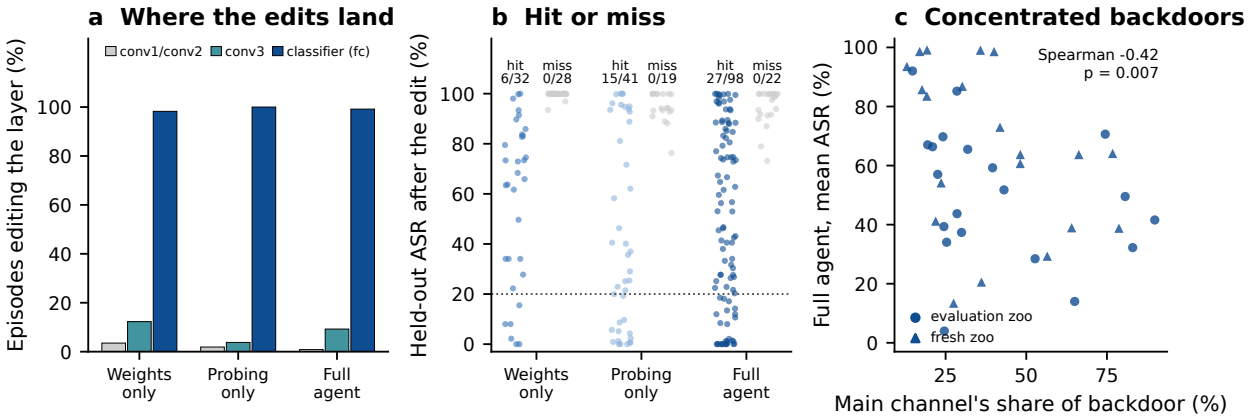


Figure 17: Anatomy of all 240 removal episodes. (a) Share of episodes whose edit changes each layer. (b) Held-out ASR after the edit, for edits that touch the network's backdoor channel ("hit", coloured) and edits that do not ("miss", grey); numbers count episodes that bring the network below 20% ASR. No edit that misses the channel removes the backdoor. (c) Share of the trigger's push towards the target carried by the backdoor channel, against the full agent's mean ASR over three runs, for the evaluation zoo (circles) and the fresh zoo (triangles).

are the ones removed (panel c): the larger the share carried by the one channel, the lower the full agent's ASR (Spearman −0.42, $p = 0.007$, 40 networks from both zoos). When the backdoor is spread over several channels, switching off the main one is not enough, and the agent rarely goes on to find the others.

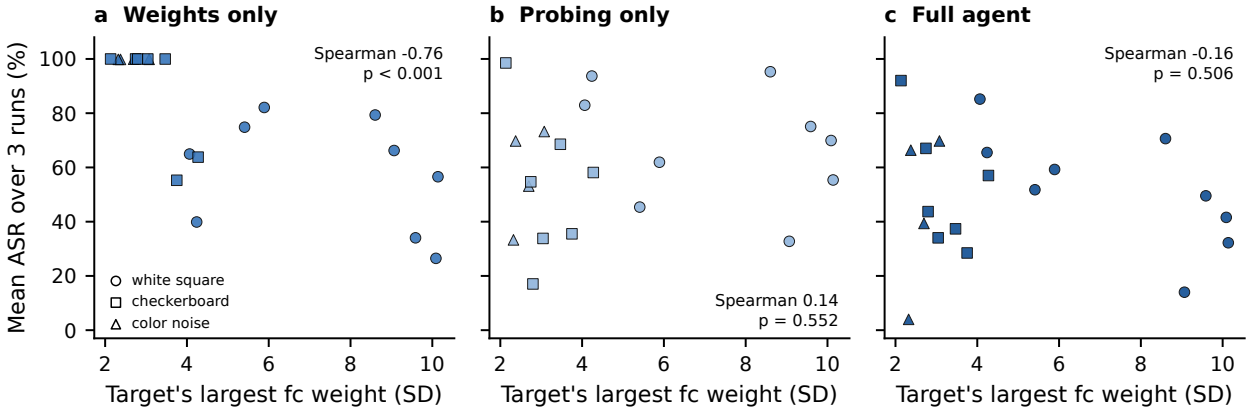


Figure 18: Reading needs a weight that stands out; probing does not. For the 20 backdoored evaluation networks, how far the target class's largest classifier weight stands out (in standard deviations of all classifier weights) against the agent's mean ASR over three runs, in each condition. Marker shape is the trigger pattern.

Figure 18 explains the ablation result of Section 6.5 at the level of single networks. Without forward passes, the agent succeeds only where the backdoor leaves an obvious mark in the weights: its ASR falls steeply as the target's largest classifier weight stands out more (Spearman −0.76, $p < 0.001$), and on all nine networks where that weight stands out by less than 3.5 standard deviations it leaves the backdoor fully intact (mean ASR at least 99.8%). With forward passes but no weight values, success does not depend on this mark at all (Spearman 0.14, $p = 0.55$): probing finds backdoors whether or not they are visible in the weights, including checkerboard and colour-noise triggers that leave little trace. The full agent sits in between (Spearman −0.16). The two routes therefore succeed on different kinds of network, which is why combining them beats either one. One confound should be noted: white-square triggers produce the most outlying weights (all networks above 5 standard deviations have white-square triggers), so for the weights-only agent "an outlying weight" and "a white square" cannot be separated with these 20 networks.

6.7 Repeat runs: how much does one run tell us?

LLM agents sample their actions, so one run per model may be luck. We repeated all three agent conditions twice more on the same 20 models with the same LLM and prompts: 120 further episodes, 180 in total (Figure 19). The averages are stable: mean ASR per run is 44.6–54.9% for the full agent, 53.5–66.0% for probing only and 73.1–80.6% for weights only, and the ordering is the same in every run. Individual models are not stable. For the full agent, a model's ASR differs by a median of 60 points between its best and worst run, and 11 of the 20 models swing by more than 50 points. The agent removes the backdoor (ASR below 20%) on 5–6 models per run, but on 11 different models across the three runs, and on only one in all three. Probing only behaves the same way (4–6 per run, 12 different models, none in all three). Weights only is steadier (median spread 2 points) because it usually changes little.

Averaging each model over its three runs gives tighter comparisons (last column of Table 7). The full agent (50.4%, CI [40.3, 59.7]) now beats probing only (60.4%, CI [50.2, 69.5]) clearly: paired difference −10.0 points, CI [−16.5, −3.4], Wilcoxon $p = 0.009$. Reading the weights therefore adds to probing, not only on the two models where it is enough on its own. Probing only beats weights only by 16.7 points ($p = 0.053$). Against NC unlearn (42.8%), the run-averaged agent is not resolved ($p = 0.35$; lower on 7 of 20 models). In 180 episodes the audit found one genuine rule break: in a repeat of the full condition, after using all 50 passes, the agent wrote its own forward pass to check clean accuracy on the 1,000 defender images. That episode left the backdoor intact (ASR 99.9%), so the results do not change whether it is counted or not. Every other flag was a comment, a shape query, or an `import sys` used to list variables.

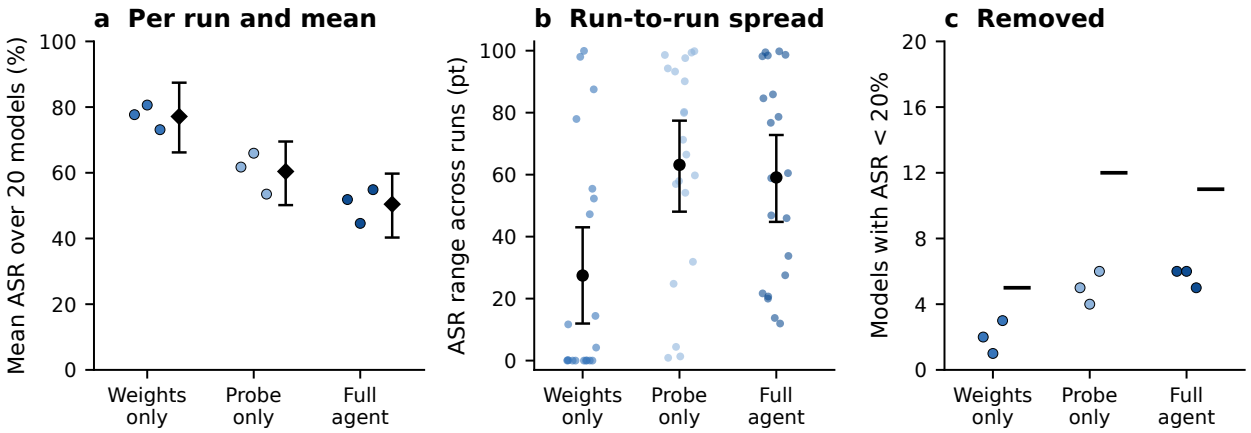


Figure 19: Run-to-run variation, three independent runs of each agent condition on the same 20 models. (a) Mean ASR of each run (dots) and the mean over runs with a 95% bootstrap CI over models (diamond). (b) Per model, ASR range across runs (best minus worst); black markers are means with 95% CIs. (c) Models brought below 20% ASR in each run (dots) and in at least one run (bar).

6.8 Best of three: can a defender pick the good run?

Since the agent often fixes a model in one run and not in another, a defender could run it several times and keep the best repair, if it can tell which one that is without the real trigger. We tested a check the defender can run, mirroring what the agent itself does: stamp 60 candidate triggers from the family the agent was told about (three patterns, sizes 3–5, four corners, three random seeds for colour noise, since the real seed is unknown) on the 1,000 defender images, and take the largest share of images any candidate sends to one class as the *estimated* ASR. We fixed the selection rule before looking at outcomes: among a model's three runs whose defender clean accuracy is at most 5 points below the original's, keep the one with the lowest estimated ASR. Selection uses no held-out data.

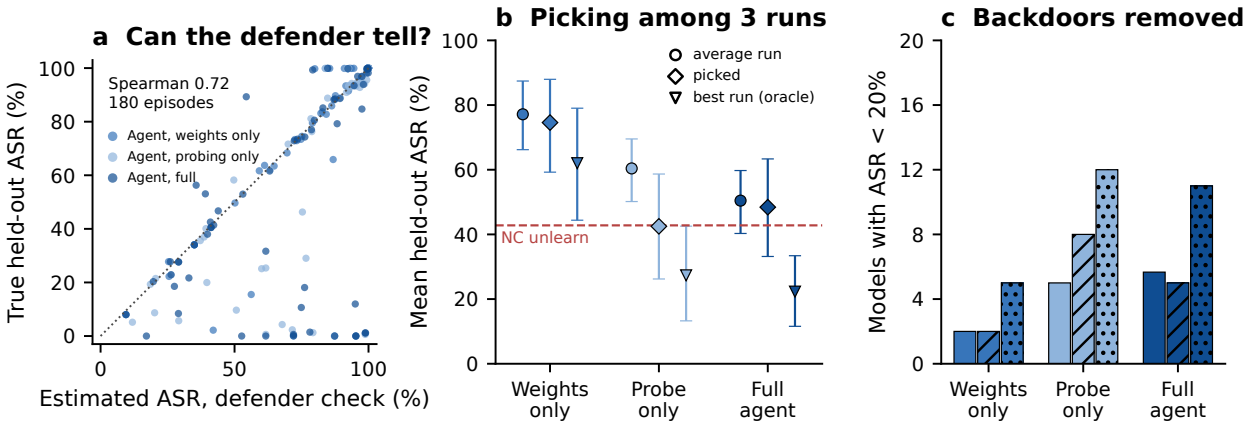


Figure 20: Picking among three runs without the real trigger. (a) Defender's estimated ASR (60 candidate triggers on its 1,000 clean images) against the true held-out ASR, for all 180 episodes. (b) Mean held-out ASR of an average run, of the run the defender picks, and of the best run (oracle); 95% bootstrap CIs over 20 models; dashed line is NC unlearn. (c) Models below 20% ASR: average per run (plain), picked (hatched), oracle (dotted).

The estimate tracks the truth reasonably well (Spearman 0.72 over 180 episodes; Figure 20a), but it errs in one direction: on some networks whose real backdoor is gone, a candidate patch still pushes many images to one class, so the check reports a high ASR. With the pre-registered rule, picking helps the probing-only agent substantially (mean ASR 60.4% for an average run, 42.5% picked, $p = 0.009$; 8 of 20 models below 20%, against 5 per run), but barely helps the full agent (50.4% to 48.4%, $p = 0.57$; 5 of 20), against an oracle of 22.3% with 11 of 20. The reason is the accuracy filter, not the estimate: in 8 of the 11 models where it missed the full agent's best run, that run cost slightly more than 5 points of defender clean accuracy and was excluded. As an *exploratory* check, chosen after seeing this, relaxing the filter to 8 points or more lowers the picked full agent to 33% mean ASR with 8 of 20 models below 20% ($p = 0.006$ against an average run), below NC unlearn's 42.8%, at a cost of about one point of held-out clean accuracy. That threshold needs confirming on fresh models before it can be trusted, but it suggests that much of the agent's run-to-run variance can be recovered by a cheap check that needs no knowledge of the real trigger. Section 6.9 tests this claim on fresh networks under a pre-registered protocol, and it does not hold up against NC unlearn.

6.9 Pre-registered confirmation on a fresh zoo

The 33% figure in Section 6.8 rests on an accuracy threshold chosen after seeing the outcomes, on the same 20 models that suggested it. To test it honestly we trained a new zoo and *pre-registered* the test before running any agent on it: the hypotheses, the selection rule with the 8-point threshold fixed, the statistical tests and what would count as failure were written down, committed and pushed to the public repository (`PREREGISTRATION.md`) before the first agent episode. The analysis script was committed at the same time.

Design. The confirmation zoo has 40 new networks (20 backdoored), trained exactly like the evaluation zoo but from a different random seed, so trigger patterns, sizes, corners, targets and poisoning rates are new draws. Its backdoored networks have 69.5% clean accuracy and 96.3% mean ASR, close to the evaluation zoo (71.2% and 96.7%). The four classical removal defences were run with unchanged settings. The full agent (same LLM, same prompt, 50 forward passes) was run three times on each of the 20 backdoored networks, 60 episodes in all, and the defender picked one run per network with the fixed rule: among runs whose defender clean accuracy is at most 8 points below the original's, the one with the lowest estimated ASR from the 60 candidate triggers. The primary hypothesis H1 was that the picked run has lower held-out ASR than NC unlearn, tested one-sided with a paired Wilcoxon signed-rank test at $\alpha = 0.05$; the secondary hypothesis H2

was that the picked run beats an average run.

Table 8: Pre-registered confirmation zoo, 20 new backdoored networks. Means with 95% bootstrap CIs on held-out images. "Picked" applies the selection rule fixed before any agent run; "oracle" takes the best of the three runs using the true ASR and is not available to a defender.

Defence	Clean accuracy (%)	Attack success (%)	ASR < 20%
No defence	69.5 [68.1, 71.1]	96.3 [94.7, 98.1]	0/20
Prune outlier	68.6 [66.9, 70.3]	82.3 [72.3, 90.9]	0/20
Fine-tune	67.3 [65.9, 68.8]	93.9 [91.0, 96.7]	0/20
Fine-pruning	66.7 [65.3, 68.1]	68.5 [53.8, 81.9]	3/20
NC unlearn	67.1 [65.6, 68.5]	47.0 [35.9, 57.6]	3/20
Agent, run 1	67.8 [66.0, 69.7]	70.7 [56.6, 82.9]	2/20
Agent, run 2	67.7 [65.9, 69.4]	61.7 [44.7, 77.5]	5/20
Agent, run 3	68.1 [66.1, 70.1]	63.6 [48.0, 78.2]	3/20
Agent, picked (8 pt, pre-registered)	67.4 [65.5, 69.3]	51.5 [36.4, 66.0]	5/20
Agent, best run (oracle)	67.1 [65.1, 68.9]	46.3 [29.7, 63.0]	7/20

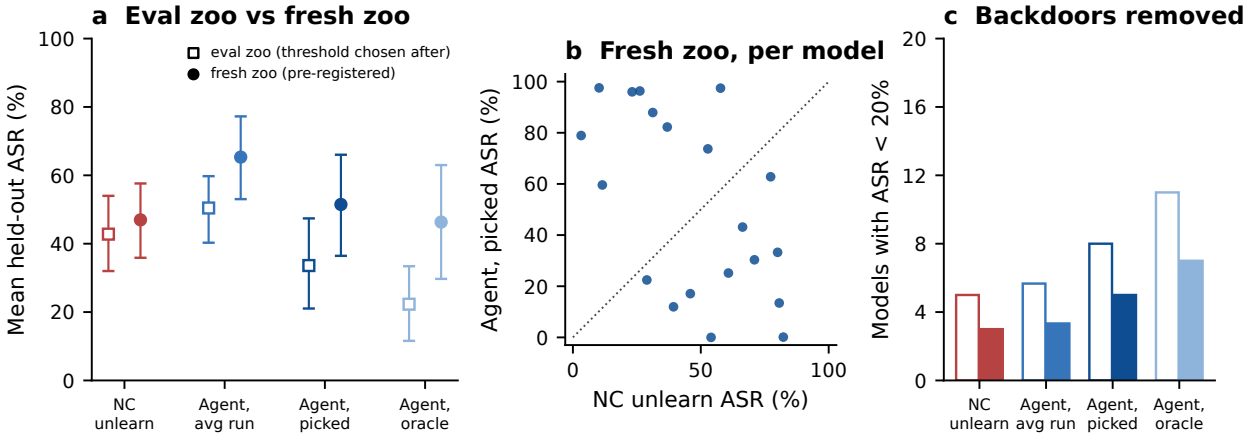


Figure 21: The best-of-three result does not replicate against NC unlearn. (a) Mean held-out ASR with 95% bootstrap CIs on the evaluation zoo (hollow; threshold chosen after seeing the results) and on the fresh zoo (filled; threshold fixed in advance). (b) Fresh zoo, per network: picked agent run against NC unlearn; the agent is lower on 11 of 20, but its losses are larger than its wins. (c) Networks brought below 20% ASR (hollow: evaluation zoo; filled: fresh zoo).

Results. Table 8 and Figure 21 give the outcome. **H1 is not supported.** The picked agent run leaves 51.5% mean ASR (CI [36.4, 66.0]) against 47.0% for NC unlearn (CI [35.9, 57.6]); the paired difference is +4.5 points (CI [−17.7, 26.9]), and the one-sided test gives $p = 0.69$. Per network the agent is lower on 11 of 20, but where it wins it wins by 39 points on average and where it loses it loses by 57: a network it fails on is often left almost fully backdoored, while NC unlearn usually reduces the backdoor partly (see below). **H2 is supported**: picking lowers the agent's mean ASR from 65.3% for an average run to 51.5% ($p = 0.0003$), and raises the number of networks below 20% from 3.3 per run on average to 5, more than any classical defence on this zoo (NC unlearn and fine-pruning remove 3 each). The oracle best run reaches 46.3% and 7 networks, so the selection rule recovers most of what can be recovered from three runs.

Why it did not replicate. Two things changed between the zoos, neither under our control. First, the agent itself did worse on the fresh networks: an average run left 65.3% ASR against 50.4% on the evaluation

zoo (Mann–Whitney $p = 0.10$, descriptive), while NC unlearn moved only from 42.8% to 47.0%. With 20 networks per zoo and large per-network variation, the evaluation zoo may simply have been a favourable draw for the agent. Second, the defender's trigger check was less informative (Spearman 0.56 between estimated and true ASR, against 0.72), so selection had less to work with. The 33% on the evaluation zoo combined a favourable zoo with a threshold fitted to it; the fresh zoo removes both.

Robustness and audit. The result does not depend on the threshold: with the evaluation zoo's 5-point rule the picked agent leaves 52.1% ($p = 0.70$). The audit flagged one of the 60 episodes, and reading it showed a genuine rule break, the second in 240 removal episodes overall: after using all 50 passes, the agent wrote its own forward pass and used it on clean and patched images to check its edit. Unlike the first break, this episode succeeded (ASR 12%) and was the run picked for its network. As pre-registered, it is kept in the main result; removing it (picking among the other two runs) raises the picked agent's mean ASR to 55.2% ($p = 0.79$), which does not change the conclusion.

Why the picked agent loses on average. Figure 22 shows the reason directly, on both zoos. NC unlearn mostly weakens backdoors partly: 60–70% of networks end between 20% and 80% ASR, and only 15% keep the backdoor intact (80% or more). A single agent run leaves the backdoor intact on a third (evaluation zoo) to half (fresh zoo) of the networks, and even the picked run leaves 10–30% intact. The agent removes more backdoors outright than NC unlearn, but its misses are more often total, and mean attack success weighs those heavily.

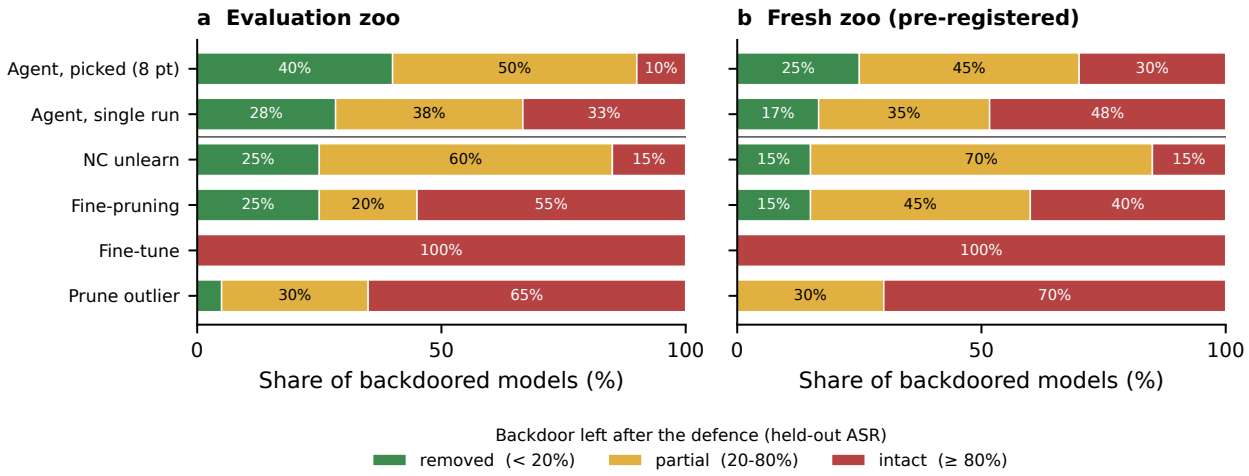


Figure 22: How each defence leaves the backdoors. Share of backdoored networks whose backdoor is removed (held-out ASR below 20%), partly removed (20–80%) or left intact (80% or more), on the evaluation zoo (a) and the fresh pre-registered zoo (b). "Agent, single run" pools the three runs (60 episodes per zoo).

7 Discussion

What the evidence supports. Table 9 lists every claim this report makes about the LLMs, the evidence for it, and how firmly it is established. Only one claim was tested under pre-registration on fresh networks, and it did not hold; the others rest on the evaluation zoo, where several analyses were chosen after seeing the data.

Why reading fails and writing works. The two tasks differ in what the LLM can do with its reasoning. In detection, it must compare thousands of numbers across a 64,000-token prompt without computing anything:

Table 9: Claims and evidence. "Confirmed": tested on fresh networks with a pre-registered protocol. "Supported": clear on the evaluation zoo (with three runs where relevant), not pre-registered. "Exploratory": chosen or tuned after seeing the data. "Not supported": failed its test.

Claim	Evidence	Status
LLMs detect backdoors from weights alone only slightly better than chance	AUROC 0.61–0.65 vs 0.87–0.94 (Fig. 7)	Supported
The LLM reader finds the same signal as a one-line weight statistic	10 of 10 targets also named by the heuristic (Fig. 8)	Supported
An LLM agent can remove backdoors by editing a few weights, with no gradients	3–6 of 20 per run, median 6 weights (Fig. 12)	Supported
The agent's edits work because of where they land	random edit 89.5% vs 51.9% ASR (Fig. 16)	Supported
Hitting the backdoor channel is necessary for removal	0 of 69 misses succeed (Fig. 17)	Supported
Removal power comes mostly from probing; reading adds to it	full vs probing only −10 pt over 3 runs, $p = 0.009$ (Table 7)	Supported
Reading works only when a weight stands out; probing does not need it	Spearman −0.76 vs 0.14 (Fig. 18)	Supported
Single agent runs are unreliable per network	11 removed in some run, 1 in all three (Fig. 19)	Supported
A defender can pick a better run without the trigger	65.3% → 51.5%, $p = 0.0003$, fresh zoo (Fig. 21)	Confirmed
The picked agent beats NC unlearn on mean ASR	33% vs 42.8% (eval, threshold chosen after)	Exploratory
	51.5% vs 47.0%, $p = 0.69$ (fresh zoo)	Not supported
The agent leaves backdoors intact more often than NC unlearn	33–48% vs 15% of networks (Fig. 22)	Supported

no means, no standard deviations, no ranking. A one-line statistic does this exactly; the LLM approximates it by eye, finds it when it is large (the cases where it names the right target coincide with the heuristic's), and otherwise falls back on its prior that backdoors are common, which inflates its probabilities on clean models. In removal, the same model can compute statistics in code, run controlled experiments, and see their results. The hard part of reading, comparing many numbers, becomes a one-line computation, and what remains is forming and testing hypotheses, which the LLM does well.

Reading and probing are complementary. The ablations (Section 6.5) refine this picture. Most of the agent's removal power comes from probing: designing candidate triggers and watching the network respond, a black-box search the LLM runs well. Reading the weights still contributes: averaged over three runs, the full agent beats the probing-only agent by 10 points ($p = 0.009$). Without any forward pass the agent removed two backdoors, with the same or nearly the same edit as the full agent, and on one of them (the case study) the probing agent, which could not see that one channel stands out, missed that edit. The full agent does best, consistent with each route covering cases the other misses: the weights suggest where to look, and experiments decide. The per-network analysis (Figure 18) makes this concrete: reading succeeds only on networks whose backdoor leaves a large outlying classifier weight, while probing succeeds regardless, including on triggers that leave no such mark.

What the agent adds. The agent's repairs are targeted: it edits a median of six weights, while the gradient-based baselines change every weight in the network. This keeps clean accuracy intact when the

target neuron is purely a trigger detector, and it makes the repair explainable: each submission comes with the neuron it blames and the evidence. The agent also succeeds on networks where the best baseline fails, and vice versa. Its main weakness is consistency: the same model is repaired in one run and untouched in the next, so a practical tool would run the agent several times and keep a repair it can verify (Section 6.8). The pre-registered test (Section 6.9) shows both sides of this: keeping the best of three runs helps the agent a great deal, but its misses are more often total than NC unlearn's (Figure 22), so on mean attack success it does not overtake a classical defence that usually weakens the backdoor partly. A practical defence could also combine the agent with classical repairs, with the agent diagnosing the trigger and choosing or calling one.

Limitations. The evaluation zoo has 40 models (20 backdoored), so confidence intervals are wide; differences such as 50.4% versus 42.8% mean ASR are not statistically resolved. Detection and the image modes were run once per model. Removal was run three times per condition, which shows large run-to-run variation per model (Section 6.7); per-model results from a single run should not be over-read. Detection modes used two different LLMs, because only one reads images, so the comparison between text and image modes is confounded with the model. Seven detection calls produced no answer because the LLM exhausted its output budget while reasoning; they are counted as uninformative. All networks share one small architecture without batch normalisation, and only patch-based BadNets attacks are studied, among the easiest backdoors to find. The agent was told that each network in the removal task is backdoored. The audit of rule violations is pattern-based and cannot exclude every indirect workaround; it found two rule breaks in 240 removal episodes (one did not affect its result; the other did and is reported with and without it), and every flagged transcript was read in full. Most analyses on the evaluation zoo were exploratory and run on the same 20 networks; the one claim we tested on fresh networks under pre-registration (best-of-three beats NC unlearn) did not replicate, and other exploratory results should be read with that in mind.

8 Future work

1. **Let the reader compute.** Allow code in detection (a weights-only sandbox with no forward passes) to test whether the LLM's weak detection comes from reading without computing.
2. **Combine the agent with classical repairs.** Give the agent Neural Cleanse, fine-pruning and unlearning as tools and let it choose from its own diagnosis, approaching the oracle combination (27.8% mean ASR) without knowing the ASR.
3. **Vary the budget and the models.** The ablations fix the two ends (0 and 50 forward passes); intermediate budgets (5, 10, 20) would show how many experiments reading saves. Compare stronger frontier models with the open models used here, and pre-register each comparison on a fresh zoo, as in Section 6.9.
4. **Make failures partial.** The agent loses to NC unlearn because its misses more often leave the backdoor fully in place (Figure 22). Running a classical repair on the networks where the agent's own trigger check still finds a working trigger would combine the agent's complete removals with the baseline's partial ones; this needs its own pre-registered test.
5. **Harder attacks and larger networks.** Blended and warping triggers, clean-label poisoning, multiple targets, attacks designed to hide from weight statistics, and architectures with batch normalisation and residual connections, where single weights are harder to read.
6. **Find the rest of a spread-out backdoor.** Hitting the main backdoor channel is necessary but often not enough (Section 6.6). After a first edit, the agent could re-run its trigger search and continue on the next channel; whether it can do this within the pass budget is untested.

7. **Faithfulness of explanations.** Check whether the neuron the agent names in its summary is the one its edit hits, across all episodes.
8. **Scale the zoo** for tighter intervals and formal tests between methods.

9 Conclusion

We built a controlled testbed for auditing neural networks through their weights and used it to test whether LLMs can detect and remove backdoors. As readers, the LLMs tested here are weak: given every weight as text or images and no way to compute, they detect backdoors only slightly better than chance and find the target class only when a simple weight statistic does. As writers, they are competitive: an LLM agent with a Python sandbox, a small forward-pass budget and no gradients finds the backdoor neuron by its own experiments, removes the backdoor with a handful of weight edits, and succeeds on networks where the best classical defence fails. Ablations show that most of this comes from probing the network, but reading the weights adds to it: over three repeated runs, the full agent beats a probing-only agent by 10 points of ASR, and reading alone solves cases that probing misses. Individual runs vary a great deal, so repeated runs, not single episodes, are the right unit of evaluation for such agents, and a simple trigger check lets a defender recover part of that variance by keeping the best of several runs. A pre-registered test on fresh networks confirmed that picking helps, but not that the picked agent beats the best classical defence on average: it removes more backdoors outright, and leaves the rest untouched. Weight-space auditing by LLMs works best when the LLM can test its hypotheses, not when it must judge by reading alone.

A Reproducibility

All code, the trained zoo, every LLM answer and transcript, and every figure are available at https://github.com/yukteswar23/backdoor-ws. Figures and tables are generated from saved results and can be rebuilt without retraining or new API calls.

Step	Command
Build the dataset file	`python -m wsp_vision.prepare_data`
Train both zoos (resumable)	`bash run_zoo.sh`
Baselines, then figures	`bash run_baselines.sh`
LLM experiments, audit, scoring	`NVIDIA_API_KEY=... bash run_nvidia.sh`
Removal ablations	`NVIDIA_API_KEY=... bash run_ablation.sh`
Repeat runs (2 and 3)	`NVIDIA_API_KEY=... bash run_repeats.sh`
Best-of-3 selection (no API)	`python -m wsp_vision.agent.select_runs`
Confirmation zoo	`python -m wsp_vision.train_zoo -split confirm -n 40`
Baselines on it	`python -m wsp_vision.baselines remove -split confirm`
Pre-registered confirmation	`NVIDIA_API_KEY=... bash run_confirm.sh`
Figures and diagrams	`python -m wsp_vision.plots.{figures,llm_figures,diagrams}`
This PDF (tables included)	`bash report/build.sh`

Compute. One NVIDIA GeForce RTX 2050 laptop GPU. Training takes about 45 s per model (about 1.5 h for 120 models); Neural Cleanse takes about 97 s per model; the full baseline suite took 1.6 h. The LLM runs used NVIDIA's free API catalog, which is rate-limited (in our measurement, a burst of about 30 requests was accepted within seconds before further requests were refused); with three removal episodes and up to seven detection calls in parallel, all LLM experiments finished in under three hours. The two LLM ablations took a further 1.5 h (probing only about 2.4 min and weights only about 10 min per model, median), with scoring on the CPU. The two repeat runs (120 episodes) took 6.3 h, slowed by API overload (HTTP 503) during the second.

B Prompt summary

Detection (system prompt, abridged): the LLM is told it audits small CIFAR-10 classifiers for hidden backdoors using only their weights; it receives the architecture, normalisation and classes; it is told that about half the networks are backdoored with a 3–5 pixel white-square, checkerboard or colour-noise patch one pixel from a corner, relabelled to one target, and that training settings were randomised identically for both groups. It must give a calibrated probability, a target class (always), a pattern and corner (always, saying in the evidence if they cannot be read), and short evidence. **Removal** (system prompt, abridged): the LLM is told that the network *is* backdoored, the trigger family, the sandbox contents and rules listed in Section 4.2, that it is scored on held-out images by attack success and clean accuracy, and that it should work step by step, call `submit(W)`, and end with a summary of what it changed and why. The full prompts are in `wsp_vision/agent/detect.py` and `remove.py`.

References

[1] S. Bills, N. Cammarata, D. Mossing, H. Tillman, L. Gao, G. Goh, I. Sutskever, J. Leike, J. Wu, and W. Saunders. Language models can explain neurons in language models. OpenAI, 2023.

[2] X. Chen, C. Liu, B. Li, K. Lu, and D. Song. Targeted backdoor attacks on deep learning systems using data poisoning. *arXiv preprint arXiv:1712.05526*, 2017.

[3] B. Efron. Bootstrap methods: another look at the jackknife. *The Annals of Statistics*, 7(1):1–26, 1979.

[4] T. Gu, B. Dolan-Gavitt, and S. Garg. BadNets: Identifying vulnerabilities in the machine learning model supply chain. *arXiv preprint arXiv:1708.06733*, 2017.

[5] A. Krizhevsky. Learning multiple layers of features from tiny images. Technical report, University of Toronto, 2009.

[6] K. Liu, B. Dolan-Gavitt, and S. Garg. Fine-pruning: Defending against backdooring attacks on deep neural networks. In *Research in Attacks, Intrusions, and Defenses (RAID)*, 2018.

[7] A. Navon, A. Shamsian, I. Achituve, E. Fetaya, G. Chechik, and H. Maron. Equivariant architectures for learning in deep weight spaces. In *International Conference on Machine Learning (ICML)*, 2023.

[8] NVIDIA. NVIDIA API catalog. https://build.nvidia.com, accessed November 2026.

[9] T. Schick, J. Dwivedi-Yu, R. Dessì, R. Raileanu, M. Lomeli, L. Zettlemoyer, N. Cancedda, and T. Scialom. Toolformer: Language models can teach themselves to use tools. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.

[10] K. Schürholt, D. Taskiran, B. Knyazev, X. Giró-i-Nieto, and D. Borth. Model zoos: A dataset of diverse populations of neural network models. In *NeurIPS Datasets and Benchmarks Track*, 2022.

[11] B. Tran, J. Li, and A. Madry. Spectral signatures in backdoor attacks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2018.

[12] T. Unterthiner, D. Keysers, S. Gelly, O. Bousquet, and I. Tolstikhin. Predicting neural network accuracy from weights. *arXiv preprint arXiv:2002.11448*, 2020.

[13] B. Wang, Y. Yao, S. Shan, H. Li, B. Viswanath, H. Zheng, and B. Y. Zhao. Neural Cleanse: Identifying and mitigating backdoor attacks in neural networks. In *IEEE Symposium on Security and Privacy (S&P)*, 2019.

[14] X. Xu, Q. Wang, H. Li, N. Borisov, C. A. Gunter, and B. Li. Detecting AI trojans using meta neural analysis. In *IEEE Symposium on Security and Privacy (S&P)*, 2021.

[15] S. Yao, J. Zhao, D. Yu, N. Du, I. Shafran, K. Narasimhan, and Y. Cao. ReAct: Synergizing reasoning and acting in language models. In *International Conference on Learning Representations (ICLR)*, 2023.